

CS 357 - 16 Optimization

Boyang Li (boyangl3)

2 types of optimization: Function $f(\mathbf{x}) : S \rightarrow \mathbb{R}$, $S \subset \mathbb{R}^n$, the point \mathbf{x}^* is called the **minimizer** or **minimum** of f if $f(\mathbf{x}^*) \leq f(\mathbf{x}) \forall \mathbf{x} \in S$. Generally, there are 2 types of optimization:

1. **Unconstrained:** Find any \mathbf{x}^* to minimize $f(\mathbf{x})$.
2. **Constrained:** Find any \mathbf{x}^* to minimize $f(\mathbf{x})$, and
 - (a) $g(\mathbf{x}) = 0$ (Equality Constraints), or
 - (b) $g(\mathbf{x}) \leq 0$ (Inequality Constraints)

While g is another function.

Introduction to maximizer: In order to find the maximizer \mathbf{x}^* for function $f(\mathbf{x})$, we can find the minimizer of function $-f(\mathbf{x})$.

Local and global minima:

- **Local minima:** \mathbf{x}^* is a local minimum if $f(\mathbf{x}^*) \leq f(\mathbf{x})$ in **some subset of domain**.
- **Global minima:** \mathbf{x}^* is a global minimum if $f(\mathbf{x}^*) \leq f(\mathbf{x})$ for **all \mathbf{x} in the domain**.

Properties for 1D stationary points:

1. **First-order derivative, necessary condition:** $f'(x^*) = 0$
2. **Second-order derivative, sufficient condition:**
 - (a) $f''(x^*) > 0$ - Local minima;
 - (b) $f''(x^*) < 0$ - Local maxima;
 - (c) $f''(x^*) = 0$ - Inflection point.

Unimodal functions:

- **Definition:** A function $f : \mathbb{R} \rightarrow \mathbb{R}$ is unimodal on an interval $[a, b]$ if this function has a unique (global) minimum on that interval $[a, b]$.
- **Properties:** Assume $x_1, x_2, x^* \in [a, b]$ and $x_1 < x_2$,
 - ◇ $x_2 < x^* \Rightarrow f(x_1) > f(x_2)$
 - ◇ $x^* < x_1 \Rightarrow f(x_1) < f(x_2)$

Golden section search:

- **Motivation:** Reduce the domain to certain $[x_1, x_2]$, find a root of the first order derivative.
- **Iteration steps:**
 1. Starting from initial interval $[a, b]$, take $x_1 = a + (1 - \tau)(b - a)$ and $x_2 = a + \tau(b - a)$
 2. Evaluate $f(x_1)$ and $f(x_2)$
 3. Update the interval $[a, b]$
 - If $f(x_1) > f(x_2)$ our new interval would be $[x_1, b]$;
 - If $f(x_1) \leq f(x_2)$ our new interval would be $[a, x_2]$;

Note: τ is the coefficient that shrinks the interval at constant rate each iteration, in the Golden section search, $\tau = \frac{\sqrt{5} - 1}{2} \approx 0.618$, which is equal to the Golden Ratio.

- **Convergence:** Linear convergent, $C = \tau \approx 0.618$
- **Cost:** One function evaluation at each iteration because we can reuse x_1 or x_2 as an interior point.
- **“Bracket length”:** Define the “bracket length” as $h = b - a$, the “bracket length” after n iterations is $h_n = \tau^n h$.

Newton’s method (for optimization):

- **Motivation:** Instead of find the root of $f(x)$, we are finding the root of $f'(x)$.
- **Iteration steps:**
 1. $f'(x_k) + f''(x_k) \cdot h_k = 0 \rightarrow h_k = -\frac{f'(x_k)}{f''(x_k)}$ (h_k is called Newton step)
 2. $x_{k+1} = x_k + h_k = x_k - \frac{f'(x_k)}{f''(x_k)}$ (This step is called Newton update)
- **Convergence:** Quadratic convergence.
- **Cost:** Each iteration we need to evaluate 2 functions, $f'(x_k)$ and $f''(x_k)$.

Hessian matrix (2nd order derivative for N-D functions):

$$\mathbf{H}_f(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}$$

Properties of N-D stationary points:

1. **Gradient (1st order derivative), necessary condition:** $f'(\mathbf{x}^*) = \vec{0}$
2. **Hessian matrix (2nd order derivative), sufficient condition:**

$\mathbf{H}_f(\mathbf{x}^*)$	Eigenvalues of $\mathbf{H}_f(\mathbf{x}^*)$	Critical point \mathbf{x}^*
Positive definite	All positive	Minimizer
Negative definite	All negative	Maximizer
Indefinite	Indefinite	Saddle point

N-D Steepest descent:

- **Motivation:** The negative of the gradient of a differentiable function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ points downhill (i.e. towards points in the domain having lower values). In other words, given a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ at a point \mathbf{x} , the function will decrease its value in the direction of “steepest descent” $-\nabla f$. This hints us to move in the direction of $-\nabla f$ while searching for the minimum until we reach the point where $-\nabla f(\mathbf{x}) = \vec{0}$.

We know the direction we need to move to approach the minimum but we still do not know the distance we need to move in order to approach the minimum. If \mathbf{x}_k was our earlier point then we select the next guess by moving it in the direction of the negative gradient:

$$\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha(-\nabla f(\mathbf{x}_k)).$$

The next problem would be to find the α , and we use the 1-dimensional optimization algorithms to find the required α . Hence, the problem translates to:

$$\mathbf{s} = -\nabla f(\mathbf{x}_k) \min_{\alpha} (f(\mathbf{x}_k + \alpha \mathbf{s}))$$

- **Iteration steps:**

1. Evaluate steepest descent: $\mathbf{s}_k = -\nabla f(\mathbf{x}_k)$
2. Perform a line search to obtain α_k (for example, Golden Section Search):

$$\alpha_k = \operatorname{argmin}_{\alpha} f(\mathbf{x}_k + \alpha \mathbf{s}_k)$$

3. Update: $\mathbf{x}_k = \mathbf{x}_k + \alpha_k \mathbf{s}_k$

- **Convergence:** Linear convergence.
- **Side note:** $\nabla f(\mathbf{x}_{k+1})$ is orthogonal to $\nabla f(\mathbf{x}_k)$.

N-D Newton's method:

- **Motivation:** Instead of find the root of $f(\mathbf{x})$, we are finding the root of $\nabla f(\mathbf{x})$.
- **Iteration steps:**
 1. $\nabla f(\mathbf{x}_k) + \mathbf{H}(\mathbf{x}_k) \cdot \mathbf{s}_k = 0 \rightarrow \mathbf{s}_k = -\mathbf{H}(\mathbf{x}_k)^{-1} \cdot \nabla f(\mathbf{x}_k)$ (\mathbf{s}_k can also be calculated by solving the system $\mathbf{H}(\mathbf{x}_k) \cdot \mathbf{s}_k = -\nabla f(\mathbf{x}_k)$).
 2. $\mathbf{x}_{k+1} = \mathbf{x}_k + \mathbf{s}_k = \mathbf{x}_k - \mathbf{H}(\mathbf{x}_k)^{-1} \cdot \nabla f(\mathbf{x}_k)$
- **Convergence:** Quadratic convergence.
- **Cost:** The cost per iteration is $\mathcal{O}(n^3)$, the cost of evaluate Hessian matrix is $\mathcal{O}(n^2)$.
- **Drawbacks:**
 - Need to evaluate 2nd order derivative
 - Only converges locally
 - Works poorly when Hessian is nearly indefinite

SUMMARY: Computational cost (only CORRECT statements):

- ✓ Secant method reuses result from previous iterations to save computational resources.
- ✓ Secant method and bisection method have similar computational cost per iteration after several iterations.
- ✓ As the number of iteration increases, Newton's method is the most computationally expensive one among the three methods.

SUMMARY: Convergence (only CORRECT statements):

- ✓ Secant method has superlinear convergence and has a lower cost for each iteration compared to Newton's methods.
- ✓ Bisection method has linear convergence if $f(x)$ is continuous within an interval $[a, b]$ such that $f(a)f(b) < 0$.
- ✓ Newton's method has quadratic convergence when it is close to the root.
- ✓ Of all three methods, Newton's method has the fastest convergence.

SUMMARY: Newton's method (1-D, N-D, find root, optimization):

	Root/solution finding	Optimization
1D	$x_{k+1} = x_k - \frac{f(x)}{f'(x)}$	$x_{k+1} = x_k - \frac{f'(x)}{f''(x)}$
N-D	$\mathbf{x}_{k+1} = \mathbf{x}_k - \mathbb{J}(\mathbf{x}_k)^{-1} \cdot f(\mathbf{x}_k)$	$\mathbf{x}_{k+1} = \mathbf{x}_k - \mathbf{H}(\mathbf{x}_k)^{-1} \cdot \nabla f(\mathbf{x}_k)$

- **Convergence:** All Newton's methods have quadratic convergence.
- **Cost:** For each iteration, we need to evaluate 2 functions.
 - For N-D root finding, the cost of find \mathbf{s} by solving a linear system is $\mathcal{O}(n^3)$; the cost of evaluation Jacobian matrix is $\mathcal{O}(n^2)$
 - For N-D optimization, the cost per iteration is $\mathcal{O}(n^3)$, the cost of evaluate Hessian matrix is $\mathcal{O}(n^2)$.
- **Drawbacks:**
 - Function need to be differentiable
 - Expensive cost to evaluate the function and derivative
 - Only converges locally (and may converges to a maxima or inflection point)
- **Concept questions on the exam (only CORRECT statements are listed):**
 - ✓ Newton's method may fail to find the minimum if the start guess is close to a maximum.
 - ✓ Newton's method will fail if the second derivative of $f(x)$ at some point x_i is 0.
 - ✓ Newton's method will not always converge.
 - ✓ The cost of solving the nonlinear system of n equations for one iteration is $\mathcal{O}(n^3)$.
 - ✓ Newton's method for solving n nonlinear equations typically has quadratic convergence when close enough to the root.
 - ✓ Newton's method can usually achieve quadratic convergence when the start guess is near the minimum.
 - ✓ The cost of evaluation Jacobian matrix is $\mathcal{O}(n^2)$
 - ✓ The cost of calculating the Hessian matrix is $\mathcal{O}(n^2)$
 - ✓ Newton method needs to evaluate derivatives.
 - ✓ Newton's method requires two function calls per iteration.